


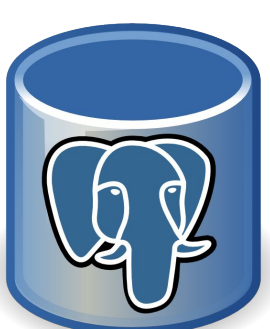
Balsa: Learning a Query Optimizer Without Expert Demonstrations

Optimizers take expert-years to develop & maintain


Optimizers are responsible for producing the best execution plan for a declarative query:




As a performance-critical component, they have been **costly to develop or maintain**:



First optimizer since pre-2000s
Commits to optimizer till occurring




Shipped first optimizer by a team
and “9 months of intense effort”




Heuristic optimizer in 2014
Cost-based opt. 3 years later


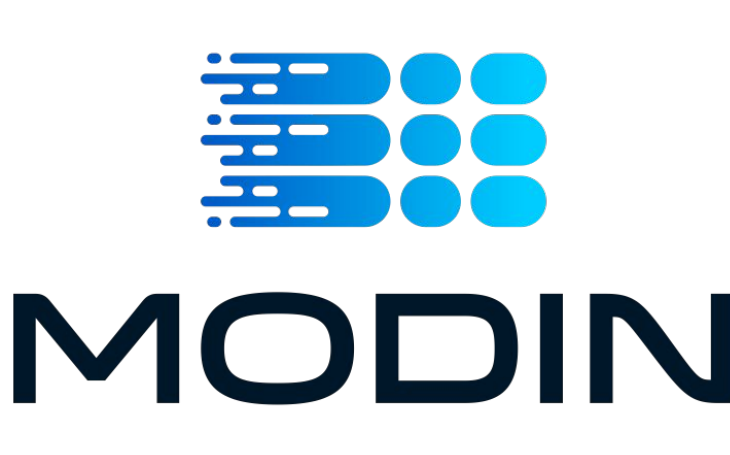

New data systems lack an “expert” to imitate from



New execution engine/model
(dynamic tasks + actors)

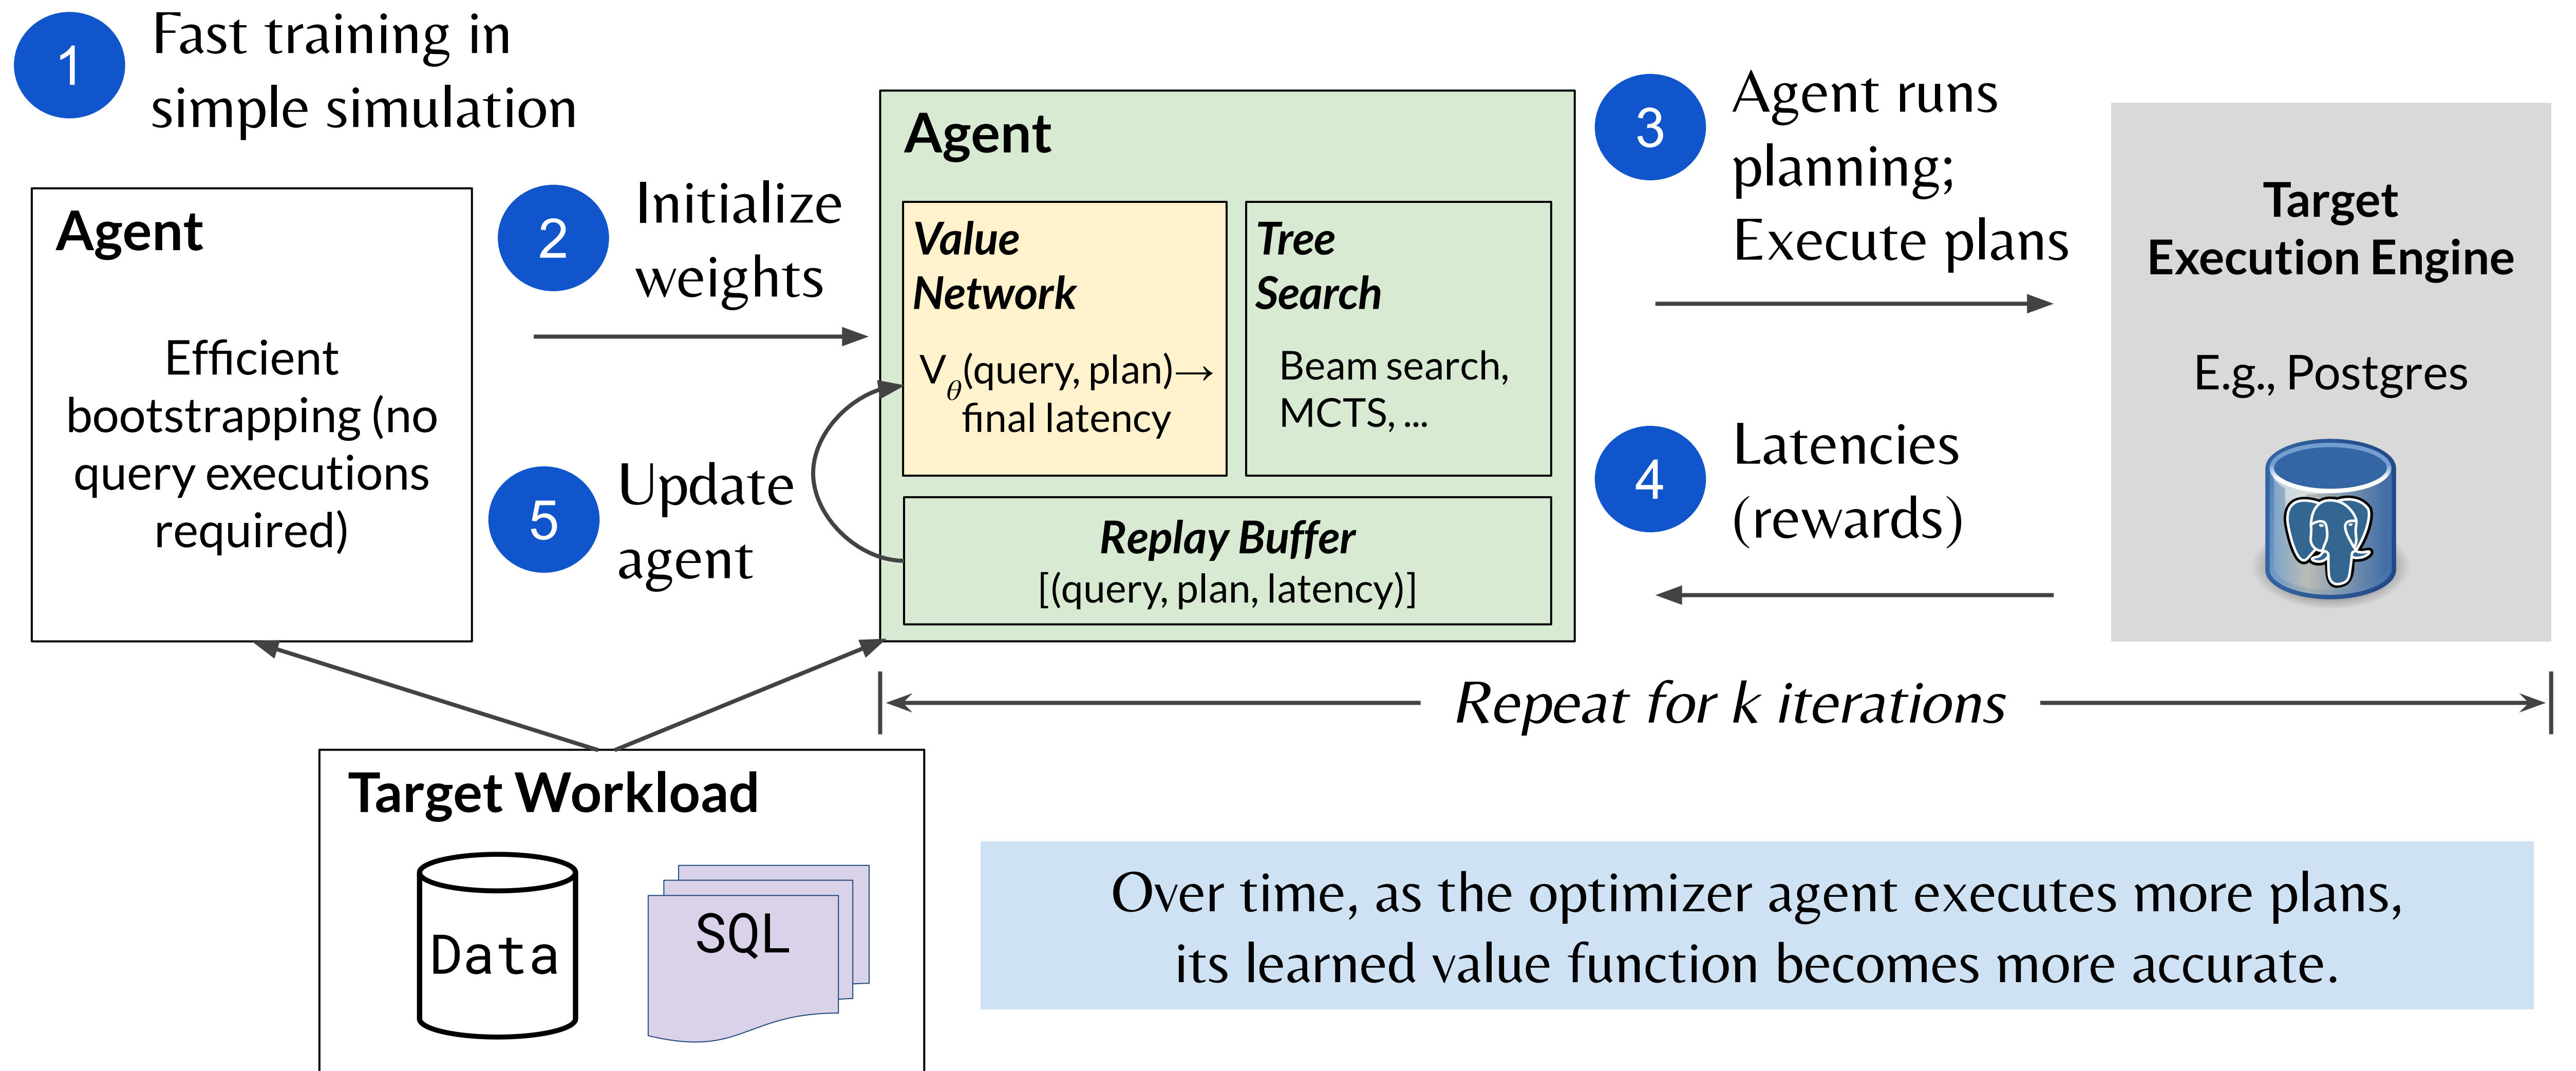


New objective:
lower memory footprint vs. latency



Balsa: a Learned Query Optimizer

Key idea: Without imitating an existing expert optimizer, learn by trial-and-error using deep RL



Key challenges

Cold start

Without imitating an expert, a cold-start agent can take forever to learn


Bad actions are slow

Many bad plans, which take too long to run (in games, bad actions speed up episodes)

Exploration

Ensuring the exponential search space is explored sufficiently

Evaluation



Takeaway: Balsa outperforms two mature expert systems & generalizes to unseen queries